



An overview of the CS-2 architecture



We have a (fairly big) box

- This contains all the cooling and infrastructure support
 - 15RU and draws around 23kW of power
- Uses standards-based power and network connections
 - 12x standard 100 Gigabit Ethernet links and converts standard TCP-IP traffic into Cerebras protocol at full line rate




- This all serves the Cerebras Wafer Scale Engine (WSE)
 - A wafer-parallel compute accelerator, containing hundreds of thousands of independent processing elements (PEs)
 - The WSE is the reason for the tutorial!




Cerebras Wafer Scale Engine (WSE)

- Physically it is about as big as a dinner plate and built upon a 7nm process technology
- The current generation CS-2 that we will use today contains:
 - Approximately 850,000 cores
 - Each core is individually programmable
 - 40GB on-chip SRAM memory
 - A total of 20PB/s aggregate memory bandwidth
- The flexibility of the individual, independent cores and the large amount of memory means that, depending on the workload, the CS-2 is capable of delivering the performance of many GPUs
 - Although of course this depends heavily on you programming it effectively, which is the topic of the tutorial today!




What types of application suit the WSE?

- Scaling poorly across multiple nodes (e.g. FFTs, particle simulation)
 - The WSE has a fabric that is high bandwidth and low-latency, allowing for excellent parallel efficiency for non-linear and highly communicative codes
 - The CS-2 system has 850k cores and can fit problems on an individual chip that take tens to hundreds of traditional small compute nodes.
- Application is constrained by data access
 - The WSE has 40 GB of SRAM uniformly distributed across the wafer that is 1 cycle away from the processing element
 - Speeds up memory access by orders of magnitude
 - The CS-2 system is capable of 1.2 Tb/s bandwidth onto the chip




WSE conceptual high-level view

- Very many individual Processing Elements (PEs)
 - These run independent of each other (e.g. their own program counter)
- PEs are connected by 2D rectangular mesh across the chip
 - 32-bit messages (called wavelets) can be communicated with neighbours in a single cycle
- The 40GB of WSE memory is distributed amongst the PEs
 - Each PE has its own private chunk of memory



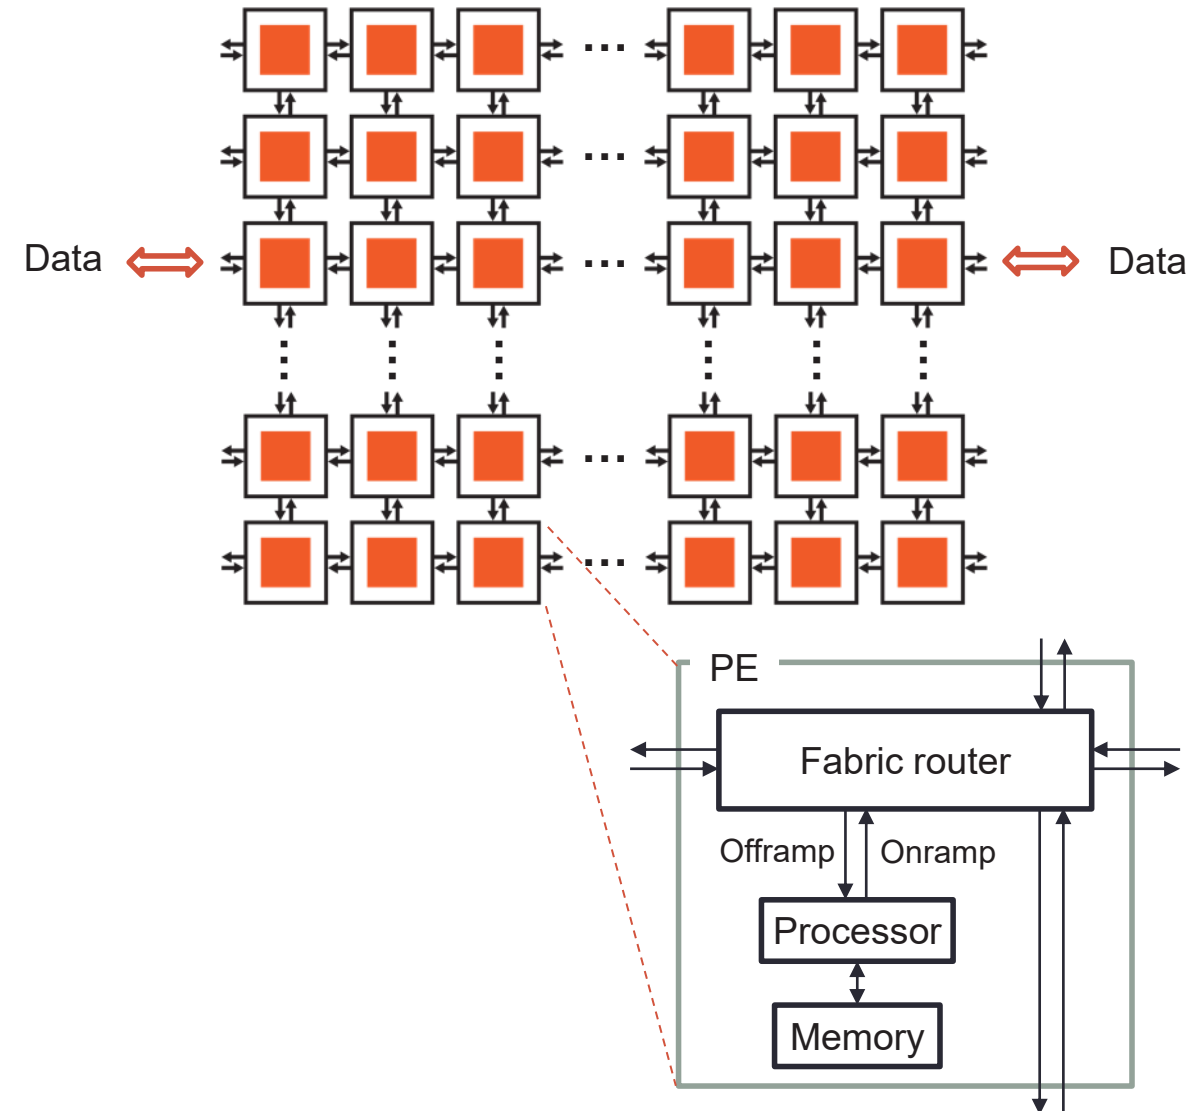
Within each Processing Element (PE)

- We have the processor itself
 - Commonly referred to as the Compute Engine (CE)
 - Independent and private from any other
- A router
 - Connected with bidirectional links to own CE and router of four neighbours
 - Link to own CE is called *the RAMP* and neighbours are referred to by north, south, east and west
 - This is the only way in which PEs can communicate
- Local (private) memory
 - All data and code for the PE is stored in this memory
 - 48KB per PE



What is supported by the CE

- 16- and 32-bit native FP and integer data types are supported by the CE
- Follows a dataflow execution model
 - Tasks are either triggered or activated
 - Independent programs specified for regions of PEs
- Control flow is straightforward to reason about
 - Tasks are non-preemptive
 - Instruction to activate another task enable state-machine behavior



From a programmer's perspective

- **Host CPU(s): Python**
 - Loads program onto simulator or CS-2 system
 - Streams in/out data from one or more workers
 - Reads/writes device memory
- **Device: Cerebras Software Language (CSL)**
 - Target software simulator or CS-2
 - CSL programs run on groups of cores on the WSE, specified by programmer
 - Executes dataflow programs



Device Read/Write



Memory I/O +
Data Streams



Programming the WSE


- Cerebras Software Language (CSL) is used to actually program the device
 - As we will see, this provides both a series of abstractions for high level programming with the ability to get down into the low level details (the same level as C) if desired
- CSL program consists of tasks and functions
 - Functions can be called by the host or another function on the device
 - Tasks are started by the hardware, run until completion, and then at that point the hardware chooses another task to run
 - Can only be *activated*, can not be called by other tasks or functions. Also don't return a value
 - Each task has an associated ID that is used for tracking and management

```
var result: f32 = 0.0;

task main_task() void {
    result = 5.0;
}
```

Communication between PEs

- Remember, a wavelet is a 32-bit message communicated with a neighbour in a single cycle
- Each physical channel has 24 virtual communication channels known as **colors** that can be used for passing wavelets
- Each wavelet has associated with it a 5-bit identify which defines which channel it is communicated on
 - Determines the wavelet's routing through the fabric and its consumption
 - This is a bit like a tag in MPI point-to-point communications, and similarly many messages on one color does not block messages with a different color using the same physical link
- Wavelets are consumed by tasks on a PE where a task is registered to execute when a wavelet arrives with a specific color




Three types of task

- Data tasks
 - Which activate upon the arrival of a wavelet with a matching color, these are used to consume messages that are communicated between PEs
 - Wavelet Triggered Task (WTT)
- Local tasks
 - Explicitly activated by other tasks or functions running on the PE
- Control tasks
 - Can be either data or local tasks, or neither (we will see examples later on)

Host runtime

- A host-side runtime known as *SdkRuntime* is provided
 - *memcpy* library loads programs, launches functions, and transfers data on and off the WSE
 - Functions provided by SdkRuntime manage the data transfer to and from the host's filesystem or memory
 - The host and WSE network interfaces finally route the data into your kernel (last step is implemented on the WSE itself to connect the I/O channel entry-points, which are in fixed locations at the edges of the WSE, to each kernel, which has a variable size and location.)
- *memcpy* infrastructure uses additional PEs around the user kernel to route tensor data and also adds a small executable component to the kernel PEs. In addition to a halo around the kernel, the additional support PEs consume three columns on the West of the kernel and two columns on the East.



The CSL compiler

- Code is compiled using the *cs/c* command
- Generates a .elf file which is the executable that will be loaded onto the CS-2
- Will play with this in a few minutes

```
usage: cslc [-h] [-o OUTPUT_NAME] [--params PARAMS] [--colors COLORS] [--memcpy] [--channels CHANNELS]
           [--import-path IMPORT_PATH] [--width-west-buf WIDTH_WEST_BUF] [--width-east-buf WIDTH_EAST_BUF] [--verbose]
           csl_filename
```

Frontend for cscl-driver. Creates a directory and then calls cscl-driver which will write its output files to the created directory.

positional arguments:

csl_filename Input CSL file

optional arguments:


-h, --help	Show this help message and exit
-o OUTPUT_NAME	Output directory name (default: out)
--params PARAMS	Comma-separated list of param-to-value mappings where a mapping is a `name:value` pair where name is a string and value is an unsigned integer. The parameter list is passed on to cscl-driver as-is.
--colors COLORS	Comma-separated list of color-to-value mappings where a mapping is a `color:value` pair where color is a string and value is an unsigned integer. The parameter list is passed on to cscl-driver as-is.
--memcpy	Add memcpy support to this program
--channels CHANNELS	Number of memcpy I/O channels to use when memcpy support is compiled with this program.
--import-path IMPORT_PATH	Add the given directory to the list of directories searched for <...> paths in @import_module and @set_tile_code statements.
--width-west-buf WIDTH_WEST_BUF	Width of west buffer (default is zero, i.e. no buffer to mitigate slow input)
--width-east-buf WIDTH_EAST_BUF	Width of east buffer (default is zero, i.e. no buffer to mitigate slow output)
--verbose	Verbose output

Use of the simulator

- CS-2 machines are designed to be exclusive use
 - Routinely developing code directly on the CS-2 would be a big pain, due to contending with lots of other users at the same time!
- Cerebras also provide a software simulator of the CS-2
 - Simulates execution of your program by running it on the CPU
 - This is highly accurate and should be used during development to test your code
 - But for performance measurements and production runs obviously we need the CS-2 itself
- We will use the simulator quite heavily during this tutorial
 - But part of our hands-on exercises will be to run each of these on the CS-2 hardware itself once you have done the development

Clustering CS-2 machines

- It's possible to combine CS-2 machines together to provide a very large virtual CS-2
- We will not cover this in the tutorial today, but details are available on the Cerebras website



Conclusions



- The CS-2 is a powerful architecture with significant raw compute and distributed memory
- We have explored the general way in which the WSE is organised and the key concepts and terminology
- Now we are going to move onto focussing on how to program the machine
 - You will see how these concepts are used practically in codes